

sql-intro-4

February 21, 2018

Create the tables for this section.

```
In [1]: %load_ext sql
        # Connect to an empty SQLite database
        %sql sqlite://
```

```
Out[1]: 'Connected: None@None'
```

```
In [2]: %%sql
        DROP TABLE IF EXISTS Purchase;
        -- Create tables
        CREATE TABLE Purchase (
            Product VARCHAR(255),
            Date     DATE,
            Price    FLOAT,
            Quantity INT
        );

        -- Insert tuples
        INSERT INTO Purchase VALUES ('Bagel', '10/21', 1, 20);
        INSERT INTO Purchase VALUES ('Bagel', '10/25', 1.5, 20);
        INSERT INTO Purchase VALUES ('Banana', '10/3', 0.5, 10);
        INSERT INTO Purchase VALUES ('Banana', '10/10', 1, 10);

        SELECT * FROM Purchase;
```

Done.

Done.

Done.

1 rows affected.

1 rows affected.

1 rows affected.

Done.

```
Out[2]: [('Bagel', '10/21', 1.0, 20),
          ('Bagel', '10/25', 1.5, 20),
          ('Banana', '10/3', 0.5, 10),
          ('Banana', '10/10', 1.0, 10)]
```

0.1 Aggregation Operations

SQL support several **aggregation** operations * SUM, COUNT, MIN, MAX, AVG * Except COUNT, all aggregations apply to a single attribute

0.1.1 COUNT

Syntax

```
SELECT COUNT(column_name)
FROM   table_name
WHERE  condition;
```

Example: Find the number of purchases

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.5	20
Banana	10/3	0.5	10
Banana	10/10	1	10

```
In [3]: %%sql
        SELECT COUNT(Product)
        FROM   Purchase;
```

Done.

```
Out[3]: [(4,)]
```

- Count applies to duplicates, unless otherwise stated
- Same as COUNT(*). Why?

Example: Find the number of **different** product purchases

- Use DISTINCT

```
In [4]: %%sql
        SELECT COUNT(DISTINCT Product)
        FROM   Purchase;
```

Done.

```
Out[4]: [(2,)]
```

0.1.2 SUM

Syntax

```
SELECT SUM(column_name)
FROM   table_name
WHERE  condition;
```

Example: How many units of all products have been purchased?

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.5	20
Banana	10/3	0.5	10
Banana	10/10	1	10

```
In [5]: %%sql
        SELECT SUM(Quantity)
        FROM   Purchase;
```

Done.

```
Out[5]: [(60,)]
```

Example: How many Bagels have been purchased?

```
In [6]: %%sql
        SELECT SUM(Quantity)
        FROM   Purchase
        WHERE  Product = 'Bagel';
```

Done.

```
Out[6]: [(40,)]
```

0.2 AVG

Syntax

```
SELECT AVG(column_name)
FROM   table_name
WHERE  condition;
```

Example: What is the average sell price of Bagels?

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.5	20
Banana	10/3	0.5	10
Banana	10/10	1	10

```
In [7]: %%sql
        SELECT AVG(Price)
        FROM   Purchase
        WHERE  Product = 'Bagel';
```

Done.

```
Out[7]: [(1.25,)]
```

0.2.1 Simple Aggregations

Example: Total earnings from Bagels sold?

```
In [8]: %%sql
        SELECT SUM(Price * Quantity)
        FROM   Purchase
        WHERE  Product = 'Bagel';
```

Done.

```
Out[8]: [(50.0,)]
```

0.3 GROUP BY

Used with aggregate functions (COUNT, MAX, MIN, SUM, AVG) to group the result-set by one or more columns.

Syntax

```
SELECT  column_name(s)
FROM    table_name
WHERE   condition
GROUP BY column_name(s)
[ORDER BY column_name(s)];
```

Example: Find total sales after 10/1 per product

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.5	20
Banana	10/3	0.5	10
Banana	10/10	4 1	10

```
In [9]: %%sql
        SELECT Product, SUM(price * quantity) AS TotalSales
        FROM Purchase
        WHERE Date > '10/1'
        GROUP BY Product;
```

Done.

```
Out[9]: [('Bagel', 50.0), ('Banana', 15.0)]
```

0.3.1 Grouping and Aggregation: Semantics of the Query

1. Compute the FROM and WHERE clauses

```
In [10]: %%sql
         SELECT *
         FROM Purchase
         WHERE Date > '10/1'
```

Done.

```
Out[10]: [('Bagel', '10/21', 1.0, 20),
          ('Bagel', '10/25', 1.5, 20),
          ('Banana', '10/3', 0.5, 10),
          ('Banana', '10/10', 1.0, 10)]
```

2. Group attributes according to GROUP BY

Product	Date	Price	Quantity
Bagel	10/21/17	1	20
	10/25/17	1.5	20
Banana	10/03/17	0.5	10
	10/10/17	1	10

Caution: SQL *only* displays one row if no aggregation function is used

```
In [11]: %%sql
         SELECT *
         FROM Purchase
         WHERE Date > '10/1'
         GROUP BY Product;
```

Done.

```
Out[11]: [('Bagel', '10/25', 1.5, 20), ('Banana', '10/10', 1.0, 10)]
```

```
In [12]: %%sql
SELECT    Product, Count(Product)
FROM      Purchase
WHERE     Date > '10/1'
GROUP BY Product;
```

Done.

```
Out[12]: [('Bagel', 2), ('Banana', 2)]
```

3. Compute the SELECT clause: grouped attributes and aggregates

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.5	20
Banana	10/3	0.5	10
Banana	10/10	1	10

```
In [13]: %%sql -- Find total sales after '10/1' per product
```

```
SELECT    Product, SUM(price * quantity) AS TotalSales
FROM      Purchase
WHERE     Date > '10/1'
GROUP BY Product;
```

Done.

```
Out[13]: [('Bagel', 50.0), ('Banana', 15.0)]
```

0.3.2 GROUP BY vs Nested Queries

```
SELECT    Product, SUM(price * quantity) AS TotalSales
FROM      Purchase
WHERE     Date > '10/1'
GROUP BY Product;
```

```
In [14]: %%sql
SELECT DISTINCT x.Product, (SELECT Sum(y.price*y.quantity)
                             FROM Purchase y
                             WHERE x.product = y.product
                             AND y.date > '10/1') AS TotalSales
FROM      Purchase x
WHERE     x.date > '10/1';
```

Done.

```
Out[14]: [('Bagel', 50.0), ('Banana', 15.0)]
```

0.4 HAVING

- HAVING clauses contain conditions on **aggregates**
- WHERE clauses condition on **individual tuples**

Syntax

```
SELECT    column_name(s)
FROM      table_name
WHERE     condition
GROUP BY  column_name(s)
HAVING    condition
[ORDER BY column_name(s)];
```

Example: Same query as before, except that we consider only products with more than 30 units sold

```
In [15]: %%sql
        SELECT    Product, SUM(price * quantity) AS TotalSales
        FROM      Purchase
        WHERE     Date > '10/1'
        GROUP BY  Product
        HAVING    SUM(Quantity) > 30;
```

Done.

```
Out[15]: [('Bagel', 50.0)]
```

1 Advanced* Topics

In this section * Relational Division in SQL * Nulls (revisited) * Outer Joins

1.1 Relational Division in SQL

- Not supported as a primitive operator, but useful for expressing queries like:

"Find suppliers who sell the x parts..."

"Find buyers who bought all products from a given category..."

- Let A have 2 fields, x and y , B have only field y

$A(x, y)$
 $B(y)$

- A/B contains all x tuples such that for every y tuple in B , there is an xy tuple in A

- Or: If the set of y values associated with an x value in A contains all y values in B , the x value is in A/B .

Classic Option 1

```
SELECT T1.x
FROM A AS T1
WHERE NOT EXISTS( SELECT T2.y
                  FROM B AS T2
                  EXCEPT
                  SELECT T3.y
                  FROM A AS T3
                  WHERE T3.y=T1.y);
```

Classic Option 2 (without EXCEPT)

```
SELECT DISTINCT T1.x
FROM A AS T1
WHERE NOT EXISTS(SELECT T2.y
                 FROM B AS T2
                 WHERE NOT EXISTS (SELECT T3.x
                                   FROM A AS T3
                                   WHERE T3.x=T1.x
                                   AND T3.y=T2.y
                                   )
                 );
```

Example: Find Establishments which sell all products

```
Establishment(eid, ename)
Sells(eid, pname)
Products(pname)
```

Classic Option 2 (without EXCEPT)

```
SELECT DISTINCT E.ename
FROM Establishment AS E
WHERE NOT EXISTS (SELECT p.pname
                  FROM Products3 AS P
                  WHERE NOT EXISTS (SELECT S.eid
                                    FROM Sells AS S
                                    WHERE S.pname=P.pname
                                    AND S.eid=E.eid
                                    )
                  );
```

Classic Option 2 (without EXCEPT)


```

SELECT DISTINCT E.ename
FROM Establishment AS E
WHERE NOT EXISTS (SELECT p.pname
                  FROM Products3 AS P
                  WHERE NOT EXISTS (SELECT S.eid
                                    FROM Sells AS S
                                    WHERE S.pname=P.pname
                                    AND S.eid=e.eid
                                    )
                  );

```

- Semantics:
- *Establishment* E such that...
 - ... there is no *Product* P...
 - without a *Sells* tuple showing that E sells P

Example: Find Establishments which sell all products

Establishment		Sells		Products 1	Products 2	Products 3
eid	ename	eid	pname	pname	pname	pname
1	Carrefour	1	Wine	Bread	Bread	Wine
2	Franprix	1	Bread		Meat	Bread
3	Boulangerie	1	Meat			Meat
4	Biocoop	1	Cheese			
		2	Wine			
		2	Bread			
		3	Bread			
		4	Bread			
		4	Meat			

```

In [16]: %%sql
DROP TABLE IF EXISTS Establishment;
-- Create tables
CREATE TABLE Establishment (
    eid INT,
    ename VARCHAR);

DROP TABLE IF EXISTS Sells;
-- Create tables
CREATE TABLE Sells (
    eid INT,
    pname VARCHAR);

DROP TABLE IF EXISTS Products1;
-- Create tables
CREATE TABLE Products1 (
    pname VARCHAR);

DROP TABLE IF EXISTS Products2;
-- Create tables

```

```

CREATE TABLE Products2 (
    pname VARCHAR);

DROP TABLE IF EXISTS Products3;
-- Create tables
CREATE TABLE Products3 (
    pname VARCHAR);

-- Insert tuples
INSERT INTO Establishment VALUES (1, 'Carrefour');
INSERT INTO Establishment VALUES (2, 'Franprix');
INSERT INTO Establishment VALUES (3, 'Boulangerie');
INSERT INTO Establishment VALUES (4, 'Biocoop');
INSERT INTO Sells VALUES (1, 'Wine');
INSERT INTO Sells VALUES (1, 'Bread');
INSERT INTO Sells VALUES (1, 'Cheese');
INSERT INTO Sells VALUES (1, 'Meat');
INSERT INTO Sells VALUES (2, 'Wine');
INSERT INTO Sells VALUES (2, 'Bread');
INSERT INTO Sells VALUES (3, 'Bread');
INSERT INTO Sells VALUES (4, 'Bread');
INSERT INTO Sells VALUES (4, 'Meat');
INSERT INTO Products1 VALUES ('Bread');
INSERT INTO Products2 VALUES ('Bread');
INSERT INTO Products2 VALUES ('Meat');
INSERT INTO Products3 VALUES ('Wine');
INSERT INTO Products3 VALUES ('Bread');
INSERT INTO Products3 VALUES ('Meat');

```

Done.

Done.

Done.

Done.

Done.

Done.

Done.

Done.

Done.

Done.

Done.

1 rows affected.

1 rows affected.

1 rows affected.

1 rows affected.

1 rows affected.

1 rows affected.

1 rows affected.

1 rows affected.

```
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.  
1 rows affected.
```

```
Out[16]: []
```

```
In [17]: %%sql  
        SELECT * FROM Establishment;
```

```
Done.
```

```
Out[17]: [(1, 'Carrefour'), (2, 'Franprix'), (3, 'Boulangerie'), (4, 'Biocoop')]
```

```
In [18]: %%sql  
        SELECT * FROM Sells;
```

```
Done.
```

```
Out[18]: [(1, 'Wine'),  
          (1, 'Bread'),  
          (1, 'Cheese'),  
          (1, 'Meat'),  
          (2, 'Wine'),  
          (2, 'Bread'),  
          (3, 'Bread'),  
          (4, 'Bread'),  
          (4, 'Meat')]
```

```
In [19]: %%sql  
        SELECT * FROM Products1
```

```
Done.
```

```
Out[19]: [('Bread',)]
```

Example: Find Establishments which sell all products

Establishment	
eid	ename
1	Carrefour
2	Franprix
3	Boulangerie
4	Biocoop

Sells	
eid	pname
1	Wine
1	Bread
1	Meat
1	Cheese
2	Wine
2	Bread
3	Bread
4	Bread
4	Meat

Products 1
pname
Bread

Products 2
pname
Bread
Meat

Products 3
pname
Wine
Bread
Meat

```
In [20]: %%sql -- Change bellow to query Products[1,2,3]
SELECT DISTINCT E.ename
FROM Establishment AS E
WHERE NOT EXISTS (SELECT P.pname
                  FROM Products1 AS P
                  WHERE NOT EXISTS (SELECT S.eid
                                    FROM Sells AS S
                                    WHERE S.pname=P.pname
                                    AND S.eid=E.eid
                                    )
                  );
```

Done.

```
Out[20]: [('Carrefour',), ('Franprix',), ('Boulangerie',), ('Biocoop',)]
```

Exercise: Write the same query with EXCEPT (Classic Option 1)

```
Establishment(eid, ename)
Sells(eid, pname)
Products(pname)
```

```
SELECT T1.x
FROM A AS T1
WHERE NOT EXISTS( SELECT T2.y
                  FROM B AS T2
                  EXCEPT
                  SELECT T3.y
                  FROM A AS T3
                  WHERE T3.y=T1.y);
```

Example: Find Establishments which sell all products

Establishment		Sells		Products 1	Products 2	Products 3
eid	ename	eid	pname	pname	pname	pname
1	Carrefour	1	Wine	Bread	Bread	Wine
2	Franprix	1	Bread		Meat	Bread
3	Boulangerie	1	Meat			Meat
4	Biocoop	1	Cheese			
		2	Wine			
		2	Bread			
		3	Bread			
		4	Bread			
		4	Meat			

```
In [21]: %%sql
-- Write the same query with EXCEPT (Classic Option 1)
```

Done.

```
Out[21]: []
```

1.1.1 Yet another option

"A Simpler (and Better) SQL Approach to Relational Division"

Journal of Information Systems Education, Vol. 13(2)

1.2 Null Values

- For *numerical operations*, NULL -> NULL:
- If x is NULL then $4 \cdot (3 - x) / 7$ is still NULL
- For *boolean operations*, in SQL there are three values:

```
FALSE = 0
UNKNOWN = 0.5
TRUE = 1
```

- If x is NULL then $x = \text{'Joe'}$ is UNKNOWN

```
C1 AND C2 = min(C1, C2)
C1 OR C2 = max(C1, C2)
NOT C1 = 1 - C1
```

Example:

```
SELECT *
FROM Person
WHERE (age < 25)
AND (height > 6 AND weight > 190);
```

Won't return: - age=20 - height=NULL <-- - weight=200

Rule in SQL: include only tuples that yield TRUE (1.0)

Example: Unexpected behavior

```
SELECT *
FROM   Person
WHERE  age < 25 OR age >= 25;
```

Some tuples from *Person* are not included
Test for NULL explicitly: * x IS NULL * x IS NOT NULL

```
SELECT *
FROM   Person
WHERE  age < 25 OR age >= 25 OR age IS NULL;
```

Now it includes all tuples in *Person*

1.3 Inner Joins + NULLS = Lost data?

- By default, joins in SQL are **inner joins**

Example: Find Products (Name) and the Stores where they are sold.

Product(name, category)
Purchase(prodName, store)

Example: Find Products (Name) and the Stores where they are sold.

Product(name, category)
Purchase(prodName, store)

Syntax 1

```
SELECT Product.name, Purchase.store
FROM   Product
JOIN   Purchase ON Product.name = Purchase.prodName;
```

Syntax 2

```
SELECT Product.name, Purchase.store
FROM   Product, Purchase
WHERE  Product.name = Purchase.prodName;
```

- Both equivalent, both *inner joins*
- **However:** Products that never sold (with no Purchase tuple) will be lost!

1.4 Outer Joins

- An **outer join** returns tuples from the joined relations that don't have a corresponding tuple in the other relations
- i.e. If we join relations A and B on $a.X = b.X$, and there is an entry in A with $X=5$, but none in B with $X=5$ LEFT [OUTER] JOIN will return a tuple (**a, NULL**)

Syntax

```
SELECT column_name(s)
FROM   table1
LEFT OUTER JOIN table2 ON table1.column_name = table2.column_name;
```

In [22]: %%sql

```
-- Create tables
DROP TABLE IF EXISTS Product;
CREATE TABLE Product (
    name VARCHAR(255) PRIMARY KEY,
    category VARCHAR(255)
);

DROP TABLE IF EXISTS Purchase;
CREATE TABLE Purchase(
    prodName varchar(255),
    store varchar(255)
);

-- Insert tuples
INSERT INTO Product VALUES ('Gizmo', 'Gadget');
INSERT INTO Product VALUES ('Camera', 'Photo');
INSERT INTO Product VALUES ('OneClick', 'Photo');

INSERT INTO Purchase VALUES ('Gizmo', 'Wiz');
INSERT INTO Purchase VALUES ('Camera', 'Ritz');
INSERT INTO Purchase VALUES ('Camera', 'Wiz');
```

Done.

Done.

Done.

Done.

Done.

1 rows affected.

1 rows affected.

1 rows affected.

1 rows affected.

1 rows affected.

Out [22]: []

```
In [23]: %%sql
        SELECT *
        FROM   Product;
```

Done.

```
Out[23]: [('Gizmo', 'Gadget'), ('Camera', 'Photo'), ('OneClick', 'Photo')]
```

```
In [24]: %%sql
        SELECT *
        FROM   Purchase;
```

Done.

```
Out[24]: [('Gizmo', 'Wiz'), ('Camera', 'Ritz'), ('Camera', 'Wiz')]
```

```
In [25]: %%sql
        SELECT Product.name, Purchase.store
        FROM   Product
        LEFT OUTER JOIN Purchase
        ON Product.name = Purchase.prodName;
```

Done.

```
Out[25]: [('Camera', 'Ritz'), ('Camera', 'Wiz'), ('Gizmo', 'Wiz'), ('OneClick', None)]
```

1.5 Outer Joins

- **Left outer join**
- Include the left tuple even if there is no match
- **Right outer join**
- Include the right tuple even if there is no match
- **Full outer join**
- Include both left and right tuples even if there is no match

2 Summary

- The relational model has rigorously defined query languages that are simple and powerful.
- Several ways of expressing a given query; a query optimizer should choose the most efficient version.
- SQL is the lingua franca (common language) for accessing relational database systems.
- SQL is a rich language that handles the way data is processed *declaratively*
- Expresses the logic of a computation without describing its control flow


```
In [26]: # Modify the css style
# from IPython.core.display import HTML
# def css_styling():
#     styles = open("./style/custom.css").read()
#     return HTML(styles)
# css_styling()
```